

ad0v1h4qd

March 18, 2024

```
[1]: import numpy as np
import pandas as pd

#import visualization libraries
import matplotlib.pyplot as plt
import seaborn as sns

#import sklearn libraries
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
```

```
[2]: ## Remove warnings
import warnings
warnings.filterwarnings('ignore')
```

```
[3]: # Load the data sets

df = pd.read_csv("diabetes.csv")
pd.set_option('display.max_columns',None)
df.head()
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[4]: df.tail()
```

```
[4]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
763             10     101             76             48     180  32.9
764              2     122             70             27        0  36.8
765              5     121             72             23     112  26.2
766              1     126             60              0        0  30.1
767              1      93             70             31        0  30.4

      DiabetesPedigreeFunction  Age  Outcome
763                   0.171    63         0
764                   0.340    27         0
765                   0.245    30         0
766                   0.349    47         1
767                   0.315    23         0
```

```
[5]: # checking shape of dataset
df.shape
```

```
[5]: (768, 9)
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null   int64
1   Glucose                             768 non-null   int64
2   BloodPressure                       768 non-null   int64
3   SkinThickness                      768 non-null   int64
4   Insulin                             768 non-null   int64
5   BMI                                 768 non-null   float64
6   DiabetesPedigreeFunction            768 non-null   float64
7   Age                                 768 non-null   int64
8   Outcome                             768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[7]: # checking missing values

df.isnull().mean()*100
```

```
[7]: Pregnancies      0.0
     Glucose        0.0
     BloodPressure   0.0
```

```

SkinThickness      0.0
Insulin            0.0
BMI                0.0
DiabetesPedigreeFunction  0.0
Age                0.0
Outcome            0.0
dtype: float64

```

```
[8]: df.columns
```

```
[8]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
          'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
          dtype='object')
```

```
[9]: df.describe() # checking summary of dataset
```

```
[9]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[10]: # calculating probability
df['Outcome'].value_counts()/len(df)*100
```


```
[10]: Outcome
0    65.104167
1    34.895833
Name: count, dtype: float64
```


```
[11]: # EDA
plt.figure()
sns.barplot(x='Outcome',y='Pregnancies',data=df)
```


```


plt.show()
plt.figure()
sns.barplot(x='Outcome' , y='Glucose' , data=df)
plt.show()
plt.figure()
sns.barplot(x='Outcome', y='SkinThickness' , data=df)
plt.show()
plt.figure()
sns.barplot(x='Outcome', y='Insulin' , data=df)
plt.show()
plt.figure()
sns.barplot(x='Outcome', y='BMI' , data=df)
plt.show()
plt.figure()
sns.barplot(x='Outcome', y='DiabetesPedigreeFunction' , data=df)
plt.figure()
sns.barplot(x='Outcome', y='Age' , data=df)
plt.show()


```



















```
[12]: plt.title('Outcome in Percentage')
plt.pie(df['Outcome'].
        value_counts(), labels=['Not_Diabetic', 'Diabetic'], autopct='%1.2f%%')
plt.show()
```

Outcome in Percentage



```
[13]: plt.figure(figsize=(13,10))  
sns.heatmap(df.corr(),annot=True)
```

```
[13]: <Axes: >
```



```
[14]: # Data preprocessing
      # Replacing Nan with mean values
df["Glucose"].fillna(df["Glucose"].mean(), inplace = True)
df["BloodPressure"].fillna(df["BloodPressure"].mean(), inplace = True)
df["SkinThickness"].fillna(df["SkinThickness"].mean(), inplace = True)
df["Insulin"].fillna(df["Insulin"].mean(), inplace = True)
df["BMI"].fillna(df["BMI"].mean(), inplace = True)
df["Pregnancies"].fillna(df["Pregnancies"].mean(), inplace = True)
df["DiabetesPedigreeFunction"].fillna(df["DiabetesPedigreeFunction"].mean(),
    ↪inplace = True)
```

```
[15]: df.describe().T
```

```
[15]:
```

	count	mean	std	min	25%	\
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	
Glucose	768.0	120.894531	31.972618	0.000	99.00000	
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	

SkinThickness	768.0	20.536458	15.952218	0.000	0.00000
Insulin	768.0	79.799479	115.244002	0.000	0.00000
BMI	768.0	31.992578	7.884160	0.000	27.30000
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375
Age	768.0	33.240885	11.760232	21.000	24.00000
Outcome	768.0	0.348958	0.476951	0.000	0.00000

	50%	75%	max
Pregnancies	3.0000	6.00000	17.00
Glucose	117.0000	140.25000	199.00
BloodPressure	72.0000	80.00000	122.00
SkinThickness	23.0000	32.00000	99.00
Insulin	30.5000	127.25000	846.00
BMI	32.0000	36.60000	67.10
DiabetesPedigreeFunction	0.3725	0.62625	2.42
Age	29.0000	41.00000	81.00
Outcome	0.0000	1.00000	1.00

0.0.1 Prediction using Logistic Regression

```
[16]: #Splitting Data
X = df.drop(columns='Outcome',axis=1)
Y = df['Outcome']
```

```
[17]: #Splitting the data
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.
↪2,stratify=Y,random_state=2)
print(X.shape)
print(X_train.shape)
print(X_test.shape)
```

```
(768, 8)
(614, 8)
(154, 8)
```

```
[18]: model = LogisticRegression()
model.fit(X_train, Y_train)
y_predict = model.predict(X_test)
```

```
[19]: y_predict
```

```
[19]: array([0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
```

```
1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1],  
dtype=int64)
```

```
[20]: #confusion matrix  
from sklearn.metrics import confusion_matrix  
confusion_matrix(y_predict,Y_test)
```


```
[20]: array([[88, 26],  
        [12, 28]], dtype=int64)
```

```
[21]: from sklearn.metrics import accuracy_score ,classification_report,f1_score
```

```
[22]: Accuracy=accuracy_score(Y_test,y_predict)  
print("Accuracy Score of test data : ",Accuracy)
```

Accuracy Score of test data : 0.7532467532467533

```
[23]: #Evaluating model  
cf_matrix = confusion_matrix(y_predict,Y_test)  
sns.heatmap(cf_matrix,annot=True,fmt=".0f")  
plt.xlabel("Predicted Values")  
plt.ylabel("Actual Values")  
plt.show()
```



```
[24]: print(classification_report(Y_test,y_predict))
```

	precision	recall	f1-score	support
0	0.77	0.88	0.82	100
1	0.70	0.52	0.60	54
accuracy			0.75	154
macro avg	0.74	0.70	0.71	154
weighted avg	0.75	0.75	0.74	154

```
[25]: #Model Testing
```

```
y_predict = model.predict([[4,130,65,30,90,33.6,0.627,50]])
print(y_predict)
if y_predict==1:
    print("You are Diabetic")
else:
    print("You are not Diabetic")
```

```
[1]
You are Diabetic
```

```
[26]: #Model Testing
```

```
y_predict = model.predict([[4,110,40,20,150,30.6,0.527,30]])
print(y_predict)
if y_predict==1:
    print("You are Diabetic")
else:
    print("You are not Diabetic")
```

```
[0]
You are not Diabetic
```

0.0.2 Prediction Random Forest Classifier

```
[27]: from sklearn.ensemble import RandomForestClassifier
      rf = RandomForestClassifier()
```

```
[28]: rf.fit(X_train,Y_train)
```

```
[28]: RandomForestClassifier()
```

```
[29]: Y_pred1 = rf.predict(X_test)
      Y_pred1
```

```
[29]: array([0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0,
          1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
          1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0,
          0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
          1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
          1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0],
          dtype=int64)
```

```
[30]: confusion_matrix(Y_pred1,Y_test)
```

```
[30]: array([[87, 25],
          [13, 29]], dtype=int64)
```

```
[31]: accuracy=accuracy_score(Y_test,Y_pred1)
      print("Accuracy Score of test data : ",accuracy)
```

Accuracy Score of test data : 0.7532467532467533

```
[32]: print(classification_report(Y_test,Y_pred1))
```

	precision	recall	f1-score	support
0	0.78	0.87	0.82	100
1	0.69	0.54	0.60	54
accuracy			0.75	154
macro avg	0.73	0.70	0.71	154
weighted avg	0.75	0.75	0.74	154

```
[33]: predict1 = rf.predict([[4,110,40,20,150,30.6,0.527,30]])
      print(predict1)
      if predict1==1:
          print("You are Diabetic")
      else:
          print("You are not Diabetic")
```

[0]
You are not Diabetic
Thank You